

PROJECT 1 ASSIGNMENT

Import the necessary libraries

```
In [1]: import pandas as pd
import numpy as np
from numbers import Number
import warnings
warnings.filterwarnings('ignore')
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression
```

Load the dataset

The encoding parameter is used to handle any special characters that may be present in the file.


```
In [2]: df = pd.read_csv('AviationData.csv', encoding='latin-1')
```

```
In [3]: # Display the first few rows of the dataframe.
df.head()
```

```
Out[3]:
```

| | Event.Id | Investigation.Type | Accident.Number | Event.Date | Location | Country |
|---|----------------|--------------------|-----------------|------------|-----------------|--------------|
| 0 | 20001218X45444 | Accident | SEA87LA080 | 1948-10-24 | MOOSE CREEK, ID | United State |
| 1 | 20001218X45447 | Accident | LAX94LA336 | 1962-07-19 | BRIDGEPORT, CA | United State |
| 2 | 20061025X01555 | Accident | NYC07LA005 | 1974-08-30 | Saltville, VA | United State |
| 3 | 20001218X45448 | Accident | LAX96LA321 | 1977-06-19 | EUREKA, CA | United State |
| 4 | 20041105X01764 | Accident | CHI79FA064 | 1979-08-02 | Canton, OH | United State |

5 rows × 31 columns



```
In [4]: # Display the first few rows of the dataframe.
df.describe()
```

Out[4]:

| | Number.ofEngines | Total.Fatal.Injuries | Total.Serious.Injuries | Total.Minor.Injuries | Tot |
|-------|------------------|----------------------|------------------------|----------------------|-----|
| count | 82805.000000 | 77488.000000 | 76379.000000 | 76956.000000 | 8 |
| mean | 1.146585 | 0.647855 | 0.279881 | 0.357061 | |
| std | 0.446510 | 5.485960 | 1.544084 | 2.235625 | |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 25% | 1.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 50% | 1.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 75% | 1.000000 | 0.000000 | 0.000000 | 0.000000 | |
| max | 8.000000 | 349.000000 | 161.000000 | 380.000000 | |

In [5]:

```
# Display the last few rows of the dataframe
df.tail()
```

Out[5]:

| | Event.Id | Investigation.Type | Accident.Number | Event.Date | Location | Coun |
|-------|----------------|--------------------|-----------------|------------|---------------|----------|
| 88884 | 20221227106491 | Accident | ERA23LA093 | 2022-12-26 | Annapolis, MD | Unit Sta |
| 88885 | 20221227106494 | Accident | ERA23LA095 | 2022-12-26 | Hampton, NH | Unit Sta |
| 88886 | 20221227106497 | Accident | WPR23LA075 | 2022-12-26 | Payson, AZ | Unit Sta |
| 88887 | 20221227106498 | Accident | WPR23LA076 | 2022-12-26 | Morgan, UT | Unit Sta |
| 88888 | 20221230106513 | Accident | ERA23LA097 | 2022-12-29 | Athens, GA | Unit Sta |

5 rows × 31 columns

In [6]:

```
df.shape
```

Out[6]:

(88889, 31)

CLEANING THE DATA

CLEANING THE DATA

In [9]:

```
print(df.isnull().sum())
```

| | |
|------------------------|-------|
| Event.Id | 0 |
| Investigation.Type | 0 |
| Accident.Number | 0 |
| Event.Date | 0 |
| Location | 52 |
| Country | 226 |
| Latitude | 54507 |
| Longitude | 54516 |
| Airport.Code | 38757 |
| Airport.Name | 36185 |
| Injury.Severity | 1000 |
| Aircraft.damage | 3194 |
| Aircraft.Category | 56602 |
| Registration.Number | 1382 |
| Make | 63 |
| Model | 92 |
| Amateur.Built | 102 |
| Number.of.Engines | 6084 |
| Engine.Type | 7096 |
| FAR.Description | 56866 |
| Schedule | 76307 |
| Purpose.of.flight | 6192 |
| Air.carrier | 72241 |
| Total.Fatal.Injuries | 11401 |
| Total.Serious.Injuries | 12510 |
| Total.Minor.Injuries | 11933 |
| Total.Uninjured | 5912 |
| Weather.Condition | 4492 |
| Broad.phase.of.flight | 27165 |
| Report.Status | 6384 |
| Publication.Date | 13771 |

dtype: int64


```
In [10]: # - Handle missing values (e.g., drop rows, fill with mean/median/mode, or use im
df.fillna(method='ffill', inplace=True) # Forward-fill missing values
```

```
In [11]: df.head()
```

Out[11]:

| | Event.Id | Investigation.Type | Accident.Number | Event.Date | Location | Country |
|---|----------------|--------------------|-----------------|------------|-----------------|---------------|
| 0 | 20001218X45444 | Accident | SEA87LA080 | 1948-10-24 | MOOSE CREEK, ID | United States |
| 1 | 20001218X45447 | Accident | LAX94LA336 | 1962-07-19 | BRIDGEPORT, CA | United States |
| 2 | 20061025X01555 | Accident | NYC07LA005 | 1974-08-30 | Saltville, VA | United States |
| 3 | 20001218X45448 | Accident | LAX96LA321 | 1977-06-19 | EUREKA, CA | United States |
| 4 | 20041105X01764 | Accident | CHI79FA064 | 1979-08-02 | Canton, OH | United States |

5 rows × 31 columns



Convert 'Publication.Date' and 'Event.Date' columns to datetime objects.

Convert specific columns to integer type.

```
In [12]: # - Convert columns to appropriate data types (e.g., date, numerical)
df['Publication.Date'] = pd.to_datetime(df['Publication.Date'])
df['Total.Fatal.Injuries'] = df['Total.Fatal.Injuries'].astype(int)
df['Total.Serious.Injuries'] = df['Total.Serious.Injuries'].astype(int)
df['Total.Minor.Injuries'] = df['Total.Minor.Injuries'].astype(int)
df['Total.Uninjured'] = df['Total.Uninjured'].astype(int)
df['Event.Date'] = pd.to_datetime(df['Event.Date'])
```

```
In [13]: # Mode imputation for 'Airport.Code'
df['Airport.Code'].fillna(df['Airport.Code'].mode()[0], inplace=True)
```

```
In [14]: df.head()
```

Out[14]:

| | Event.Id | Investigation.Type | Accident.Number | Event.Date | Location | Country |
|---|----------------|--------------------|-----------------|------------|-----------------|---------------|
| 0 | 20001218X45444 | Accident | SEA87LA080 | 1948-10-24 | MOOSE CREEK, ID | United States |
| 1 | 20001218X45447 | Accident | LAX94LA336 | 1962-07-19 | BRIDGEPORT, CA | United States |
| 2 | 20061025X01555 | Accident | NYC07LA005 | 1974-08-30 | Saltville, VA | United States |
| 3 | 20001218X45448 | Accident | LAX96LA321 | 1977-06-19 | EUREKA, CA | United States |
| 4 | 20041105X01764 | Accident | CHI79FA064 | 1979-08-02 | Canton, OH | United States |

5 rows × 31 columns




In [15]:

```
df.head()
```

Out[15]:

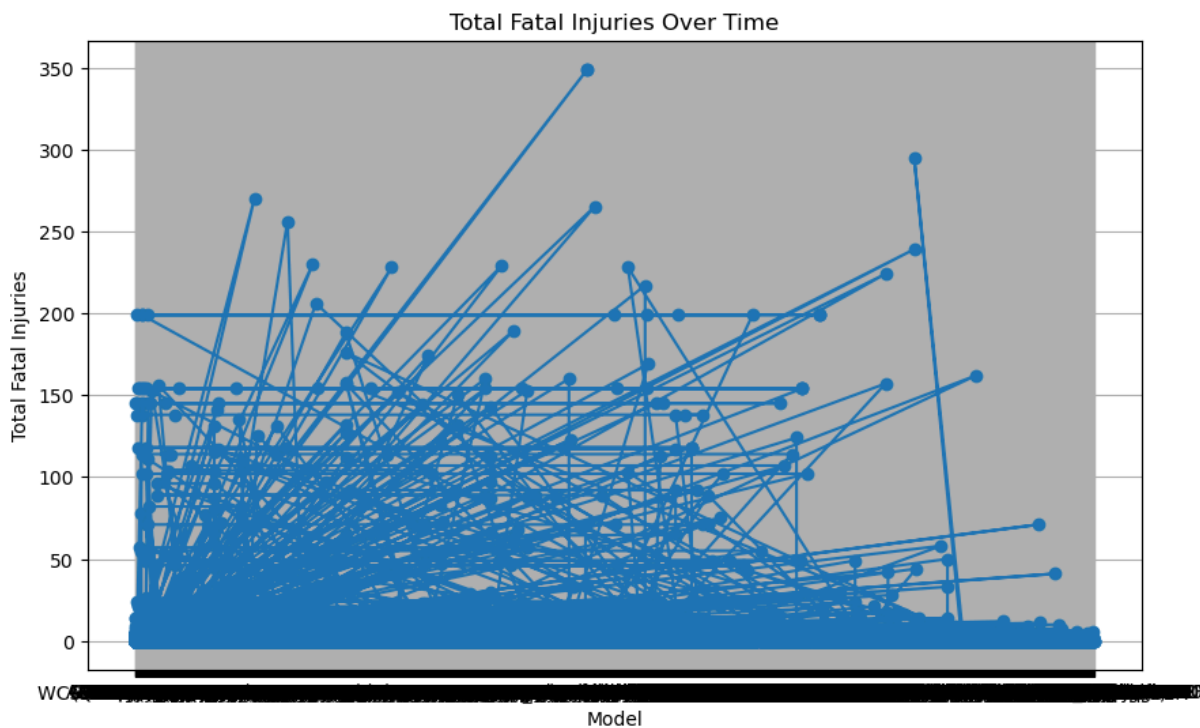
| | Event.Id | Investigation.Type | Accident.Number | Event.Date | Location | Country |
|---|----------------|--------------------|-----------------|------------|-----------------|---------------|
| 0 | 20001218X45444 | Accident | SEA87LA080 | 1948-10-24 | MOOSE CREEK, ID | United States |
| 1 | 20001218X45447 | Accident | LAX94LA336 | 1962-07-19 | BRIDGEPORT, CA | United States |
| 2 | 20061025X01555 | Accident | NYC07LA005 | 1974-08-30 | Saltville, VA | United States |
| 3 | 20001218X45448 | Accident | LAX96LA321 | 1977-06-19 | EUREKA, CA | United States |
| 4 | 20041105X01764 | Accident | CHI79FA064 | 1979-08-02 | Canton, OH | United States |

5 rows × 31 columns



In [16]:

```
# Analyze accident trends over time
plt.figure(figsize=(10, 6))
plt.plot(df['Model'], df['Total.Fatal.Injuries'], marker='o')
plt.title('Total Fatal Injuries Over Time')
plt.xlabel('Model')
plt.ylabel('Total Fatal Injuries')
plt.grid(True)
plt.show()
```



```
In [31]: # prompt: clean this data for me

# Check for and handle any remaining missing values after initial cleaning
print(df.isnull().sum())

# Further data cleaning based on specific column needs
# Example: If 'Latitude' or 'Longitude' have missing values, consider imputation or
# ... (add more specific cleaning steps as needed)

# Remove rows where 'Total.Fatal.Injuries' are greater than the total number of people
total_injured = df['Total.Fatal.Injuries'] + df['Total.Serious.Injuries'] + df['Total.Minor.Injuries']
df = df[df['Total.Fatal.Injuries'] <= total_injured]

# Ensure consistent data types across the dataframe
for col in df.columns:
    if df[col].dtype == 'object':
        try:
            df[col] = pd.to_numeric(df[col])
        except ValueError:
            pass # Skip if conversion to numeric fails

# Example: Handling inconsistent values in a categorical column (replace with appropriate values)
df['Make'] = df['Make'].replace({'CESSNA AIRCRAFT': 'CESSNA'}) # Replace inconsistent values

# Drop rows where 'Total.Fatal.Injuries' is negative (if any)
df = df[df['Total.Fatal.Injuries'] >= 0]


df.head()
```

```
Event.Id          0
Investigation.Type 0
Accident.Number   0
Event.Date        0
Location          0
Country           0
Latitude          2
Longitude         2
Airport.Code      0
Airport.Name      7
Injury.Severity   0
Aircraft.damage   0
Aircraft.Category 5
Registration.Number 0
Make             0
Model            0
Amateur.Built     0
Number.of.Engines 0
Engine.Type       0
FAR.Description   5
Schedule          5
Purpose.of.flight 0
Air.carrier       5
Total.Fatal.Injuries 0
Total.Serious.Injuries 0
Total.Minor.Injuries 0
Total.Uninjured   0
Weather.Condition 0
Broad.phase.of.flight 0
Report.Status     0
Publication.Date  1
dtype: int64
```

Out[31]:

| | Event.Id | Investigation.Type | Accident.Number | Event.Date | Location | Country |
|---|----------------|--------------------|-----------------|------------|-----------------|---------------|
| 0 | 20001218X45444 | Accident | SEA87LA080 | 1948-10-24 | MOOSE CREEK, ID | United States |
| 1 | 20001218X45447 | Accident | LAX94LA336 | 1962-07-19 | BRIDGEPORT, CA | United States |
| 2 | 20061025X01555 | Accident | NYC07LA005 | 1974-08-30 | Saltville, VA | United States |
| 3 | 20001218X45448 | Accident | LAX96LA321 | 1977-06-19 | EUREKA, CA | United States |
| 4 | 20041105X01764 | Accident | CHI79FA064 | 1979-08-02 | Canton, OH | United States |

5 rows × 31 columns



In [33]: *# prompt: Analyse for me this data. The analysis should yield three concrete business insights*
Analyse injury severity distribution

```


injury_columns = ['Total.Fatal.Injuries', 'Total.Serious.Injuries', 'Total.Minor.In
df[injury_columns].describe()


plt.figure(figsize=(12, 6))
sns.histplot(df['Total.Fatal.Injuries'], kde=True)
plt.title('Distribution of Total Fatal Injuries')
plt.xlabel('Total Fatal Injuries')
plt.ylabel('Frequency')
plt.show()


# Investigate the relationship between aircraft make and fatal injuries
plt.figure(figsize=(12, 6))
sns.boxplot(x='Make', y='Total.Fatal.Injuries', data=df)
plt.xticks(rotation=90)
plt.title('Fatal Injuries by Aircraft Make')
plt.show()

# Example: Analyze the correlation between features
correlation_matrix = df[['Total.Fatal.Injuries', 'Total.Serious.Injuries', 'Total.M
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix of Injury Types')
plt.show()

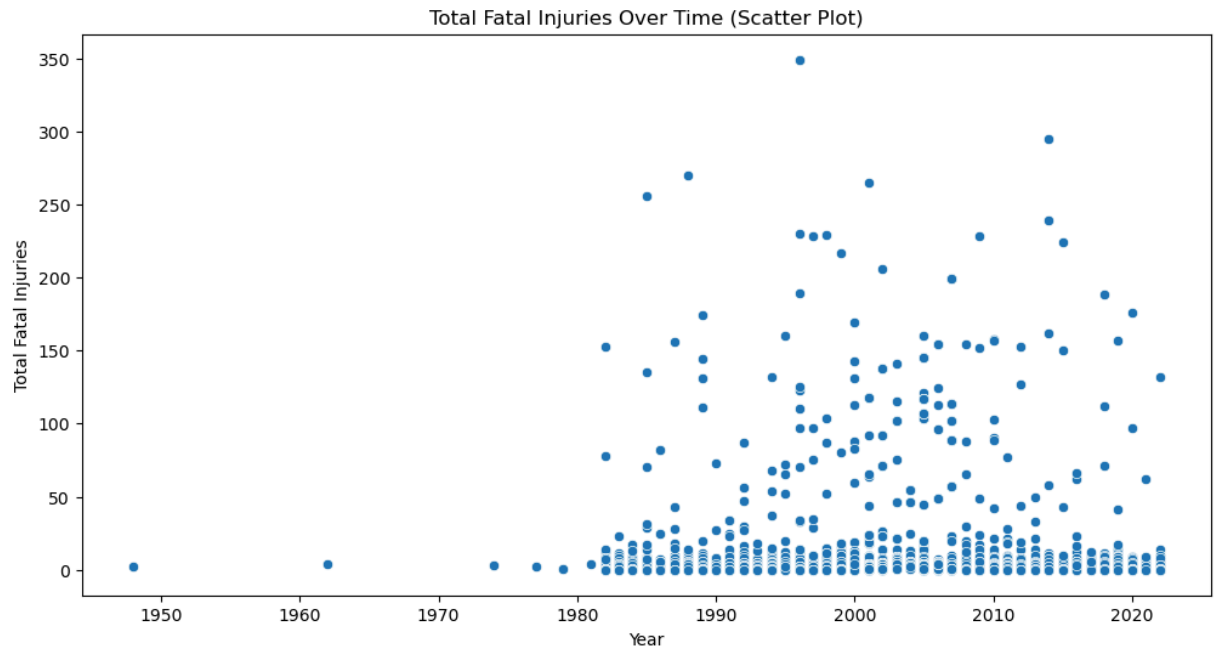
```







```
In [35]: # Analyze the relationship between the year of the event and total fatal injuries
df['Event.Year'] = df['Event.Date'].dt.year
plt.figure(figsize=(12, 6))
sns.scatterplot(x='Event.Year', y='Total.Fatal.Injuries', data=df)
plt.title('Total Fatal Injuries Over Time (Scatter Plot)')
plt.xlabel('Year')
plt.ylabel('Total Fatal Injuries')
plt.show()
```



In []: